Does hiding the post-processing from a memory-limited eavesdropper raise the noise tolerance of QKD? Timelocked advantage distillation and the 20% barrier

Hikaru Wakaura[1, *]

[1] *QIRI (Quantum Integrated Research Institute Inc.), 1–16–3 Akasaka, Minato-ku, Tokyo 107–0061, Japan*

(Dated: October 21, 2026)

Advantage distillation (AD) raises the finite-key threshold QBER of decoy-state BB84 from 9.5% to 17.3% [1]; asymptotically two-way post-processing tolerates 20%, and Bae and Acín showed that an adversary who stores her probes until the AD parities are announced attains this limit, while one forced to measure each probe beforehand is held off until the 25% entanglement limit. We ask what a quantum computational timelock on the post-processing announcements can buy for this protocol, treating the parities, the block partition and the bases separately. Hiding the parities buys nothing: we prove that for every block size the best measurement made before the parities are known attains the adaptive Helstrom value, so the key length in the min-entropy accounting is the collective-attack value, 17.66% at 10^{14} pulses against the 17.30% certified by the entropic-uncertainty bound of Ref. [1]. Hiding the partition is different: the protocol of Ref. [1] partitions the sifted key at random after the quantum phase, so an adversary whose memory is shorter than the time between the probes of a block cannot bring them together. Modelling her by role-blind per-probe measurements, an assumption we state rather than derive, the Helstrom benchmark, which a numerical search over identical measurements does not improve at odd block sizes, gives 20.8% (Shannon) and 20.7% (min-entropy) at $b = 9$, 19.4% and 20.6% in a post-selection sensitivity check, with 25% as asymptote. Whether a memory holding many probes without the partition can do better is open; at up to six probes a joint measurement recovers most of the adaptive value, so at these sizes the gain rests on locality. Hiding the basis leaves the BB84 guessing probability unchanged at one- and two-bit blocks and, for six-state, shrinks the memoryless gain by an order of magnitude at two- and three-bit blocks.

I. INTRODUCTION

Quantum key distribution promises keys whose secrecy does not rest on computational assumptions, but the promise is only redeemed when the channel is quiet enough. Every QKD protocol has a maximum QBER above which no key can be certified, and for the most deployed protocol, decoy-state BB84 with one-way post-processing, that maximum is close to 9–11% depending on the source model [1, 2]. Daylight free-space links, fibres shared with bright classical traffic, and long satellite passes can approach this level, so raising the threshold without changing the optics is one of the few remaining software levers on QKD range.

Advantage distillation is the classical lever. Introduced by Maurer for classical key agreement [3] and brought into QKD by Gottesman and Lo [4], it post-selects blocks of raw bits whose announced parities agree, so that the surviving bits are far better correlated between Alice and Bob than between Alice and Eve. Chau showed that BB84 tolerates 20.0% and six-state QKD 27.6% asymptotically with iterated two-way post-processing [5]; Acín and co-workers proved these values tight for the repetition-code protocol under collective attacks, with an explicit attack in which Eve stores her probes until the parities are public and then discriminates two product states [6, 7]. A finite-key treatment for sending-or-not-sending twin-field QKD was given by Jiang *et al.* [10], sharper finite-blocklength entropy bounds for two-way distillation have recently been obtained from hypothesis testing [9], and for decoy-state BB84 the first complete finite-key analysis is that of Treplin, Kleinpaß and Orsucci [1], who proved that a b-bit repetition-code AD step raises the finite-key threshold from 9.5% to 17.3% at $N = 10^{14}$ pulses. Their formula bounds Eve's smooth min-entropy about the distilled bits through an entropic uncertainty relation on the post-selected state, and it allows the adversary to hold a purification of everything until the last message of the protocol has been announced.

A different line of work asks what changes if the adversary cannot wait. Bae and Acín already observed that a delay between distribution and reconciliation, longer than the decoherence time of Eve's memory, forces her to measure before the parities are announced, and that for the parity-announcing form of AD this restores the entanglement limit of 25% for BB84 as the security condition, whereas for the form in which Alice announces positions of equal bits it restores nothing [7]. In the cryptographic literature the same assumption is the bounded- and noisy-storage model [11–13], under which two-party primitives become possible. Watanabe, Matsumoto and Uyematsu encrypted the parities of a two-way pre-processing step with a pre-shared one-time pad and obtained a higher net key rate in the Shor–Preskill framework [14], at the price of a pre-shared key whose provenance the analysis leaves open, as Novak later noted

* h.wakaura@qiri.co.jp

[17]. Pastushenko and Kronberg encrypted the error-correction data and argued that, if Eve's quantum memory decoheres before she learns the codeword set, her information is bounded by the additive accessible information rather than by the Holevo quantity [15]; they did not claim composable security. Novak then proved, in the non-asymptotic framework of Tomamichel and Leverrier [16], that one-time-pad protection of one-way error-correction data drawn from the same session leaves the key rate exactly unchanged at fixed security level [17], and Fung, Ma, Chau and Cai had shown that delaying privacy amplification leaves security and key rate unchanged [18]. Neither encryption nor delay by itself buys anything; whatever gain exists must come from an assumption about the adversary's memory. Vyas and Alléaume made that assumption explicit as the quantum computational hybrid model, in which a computationally secure cipher cannot be broken before a time t_{comp} and any quantum memory decoheres within $t_{\mathrm{coh}} < t_{\mathrm{comp}}$; their MUB-QCT protocol timelocks the basis information of a high-dimensional encoding and achieves everlasting security at rates beyond QKD [19, 20].

This paper asks what the same assumption can buy for the most deployed protocol when the hardware is left untouched and only post-processing announcements are timelocked, and it sharpens the question raised by the remark of Ref. [7]: what exactly must be assumed about the adversary's memory, and about which announcements are hidden, for the parity-announcing form of AD to beat 20%? Under the timelock the adversary must measure before she learns the hidden announcements, so her side information is classical and non-adaptive, and the problem becomes one of state discrimination with post-measurement information [25–27]. Our contributions are: (i) a threat model that lists the announcements of Protocol 1 of Ref. [1] with their timing, states which of them a timelock can hide, and identifies the adversary class that each choice enforces; (ii) Theorem 1: for every block size the best measurement made before the parities are announced equals the adaptive Helstrom measurement, with an explicit basis, so hiding the parities does not change the adversary's guessing probability, and the per-probe reading of the remark of Ref. [7] requires more than a hidden parity; (iii) a finite-key key length that reuses the decoy-state statistics of Ref. [1] and replaces only the adversary term, maximised over Bell-diagonal collective attacks, validated by reproducing all four published thresholds of that work, and evaluated in two accountings whose roles differ: the min-entropy accounting is the collective-attack value and does not depend on the timelock (it certifies 17.66% at $N = 10^{14}$ against 17.30% for the entropic-uncertainty bound), while the Shannon accounting is the only place where a timelock enters, and its best value exceeds the best timelock-free certificate by 0.7 percentage points for block-local measurements; (iv) for a timelock on the block partition, a model of an adversary with short memory as role-blind per-probe measurements, stated as an assumption rather than derived, in which the Helstrom-plus-tally benchmark is not improved by a numerical search over identical four-outcome measurements at odd block sizes and gives 20.7–20.8% at $N = 10^{14}$ and $b = 9$ with the 25% limit of Gisin and Wolf [21] as asymptote, together with semidefinite programmes showing that an adversary who can hold several probes recovers most of the adaptive value even without the partition, so that at these sizes the gain rests on locality; and (v) numerical evidence that the gain from hiding the basis shrinks by an order of magnitude once blocks of two or three bits are formed. We state the remaining open question in Sec. VII and record two by-products for implementations of Ref. [1].

II. MODEL AND ASSUMPTIONS

A1 (quantum computational timelock). A cipher E_k shared by Alice and Bob cannot be broken by the adversary before time t_{comp} after its use; any quantum memory available to the adversary decoheres within $t_{\mathrm{coh}} < t_{\mathrm{comp}}$ [19]. The technology assessment in Ref. [19] places t_{coh} of optical quantum memories below one second, whereas t_{comp} for a post-quantum cipher is measured in years; the assumption is only required to hold during the protocol run, which is what makes the resulting key everlasting.

A2 (announcements and their timing). Protocol 1 of Ref. [1] runs for a time T_{run} (10^5 s for $N = 10^{14}$ pulses at 1 GHz) and then announces, in this order: the detection record, the bases and intensities, a random reordering of the sifted strings, the test set, a random partition of the remaining string into blocks of length b, the AD parity strings c_A, c_B, and the reconciliation syndrome. A message m sent at time t_m may be hidden for a delay Δ_m, realised as a plain delay before publication or as encryption under E_k (then $\Delta_m = t_{\mathrm{comp}}$, counted from the publication of the ciphertext at t_m); an adversary who intercepted a probe at t_{rx} must measure it before $t_{\mathrm{rx}} + t_{\mathrm{coh}}$, so the message is hidden from that measurement iff $\Delta_m > t_{\mathrm{coh}} - (t_m - t_{\mathrm{rx}})$. Since every announcement is made after the run, an adversary with $t_{\mathrm{coh}} \ll T_{\mathrm{run}}$ measures every probe, except those received in the last t_{coh} of the run (a fraction $t_{\mathrm{coh}}/T_{\mathrm{run}}$ that we neglect), before any of them, whether or not it is additionally hidden; the timelock matters for an adversary with $t_{\mathrm{coh}} \gtrsim T_{\mathrm{run}}$, for whom it decides which announcements she sees before measuring. We therefore distinguish three adversary classes by what is known at measurement time.

A3 (physical layer). As in Ref. [1]: phase-randomised coherent states with three intensities $\mu_1 > \mu_2 + \mu_3$, $\mu_3 \approx 0$, active basis choice, threshold detectors with random assignment of double clicks, and an authenticated classical channel.

A4 (adversary classes). Each pulse interacts with an independent ancilla through a collective attack, which after the Pauli twirl of Appendix B is Bell-diagonal and, in the symmetric case, is the attack of Refs. [21, 22]. Before her memory expires the adversary may apply any

measurement to the ancillas she holds; afterwards her side information is classical. *Class J (joint, block-local).* She knows which b ancillas form a block and holds them together, but not the parities: this is the class an adversary with $t_{\mathrm{coh}} \gtrsim T_{\mathrm{run}}$ faces when only the parities and the syndrome are hidden, and it is the class that Theorem 1 settles. *Class W (window).* She does not know the partition when she measures, and each probe must be measured within t_{coh} of its arrival, so her memory holds only the probes that arrived within the last t_{coh}. For $t_{\mathrm{coh}} \ll T_{\mathrm{run}}$ the arrival times of two probes of one block, which the random partition makes uniform over the run, differ by less than t_{coh} with probability $2r - r^2$, $r = t_{\mathrm{coh}}/T_{\mathrm{run}}$, so a union bound over pairs gives a fraction of at most $2\binom{b}{2}r$ of the blocks with two probes in one memory (7×10^{-4} for $b = 9$, $t_{\mathrm{coh}} = 1$ s, $T_{\mathrm{run}} = 10^5$ s; this bound is not conditioned on sifting and acceptance), and for a block whose probes are never held together the measurement record factorises into a per-probe instrument that may depend on the outcomes of other probes but not on which of them are its block mates. Up to the feed-forward on unrelated probes, which we neglect, and up to correlations between the records of different blocks that a window measurement spanning many blocks could create, which we also neglect so that the records of distinct blocks are treated as independent, this is the class of role-blind per-probe measurements: identical or randomised single-probe POVMs followed by classical decoding once the partition and the parities are announced. Both neglects are assumptions of the model, not consequences of it. *Class H (Helstrom benchmark).* Each ancilla is measured by the Helstrom measurement of its two possible states and the outcomes are decoded optimally; this is the model of Ref. [21] and the "individual attack" of Ref. [7]. Class H is a sub-class of W; in Sec. III a numerical search over identical four-outcome per-probe measurements finds no improvement on it for odd $b \le 9$ and a strict improvement for even b.

A5 (statistics). Finite-size fluctuations of the decoy estimates, the block post-selection and the error-rate extrapolation are handled exactly as in Ref. [1]; we also verify the headline numbers with the virtual-i.i.d. reduction of Ref. [10].

The role of the cipher follows from A2. A plain delay $\tau_{\mathrm{ann}} > t_{\mathrm{coh}}$ is sufficient and needs no cryptography, at the price of a key-delivery latency of τ_{ann}. Encryption under a post-quantum cipher removes that latency while keeping the adversary's delay at t_{comp}. The cipher compresses latency; it does not create the gain, which is the reading of Ref. [15] that survives the equivalence theorem of Ref. [17].

III. THE ADVERSARY'S INFORMATION ON A DISTILLED BIT

Under A1 and A2 the adversary's optimal strategy is an immediate measurement followed by classical decoding once the delayed messages become available (Ref. [19], Proposition 1). In the symmetric attack each pulse interacts as

$$U|\uparrow\rangle|0\rangle = \sqrt{F}|\uparrow\rangle\psi_{\uparrow} + \sqrt{D}|\downarrow\rangle\varphi_{\uparrow}, \tag{1}$$

$$U|\downarrow\rangle|0\rangle = \sqrt{F}|\downarrow\rangle\psi_{\downarrow} + \sqrt{D}|\uparrow\rangle\varphi_{\downarrow}, \tag{2}$$

(Eqs. (1)–(2) of Ref. [21]), with disturbance D equal to the single-photon QBER and, by the symmetry of the attack, $F = 1 - D = (1 + c_\varphi)/(2 - c_\psi + c_\varphi)$ (Eq. (3) of Ref. [21]), where $c_\psi = \langle\psi_\uparrow|\psi_\downarrow\rangle$ and $c_\varphi = \langle\varphi_\uparrow|\varphi_\downarrow\rangle$ are real. The ψ and φ states span orthogonal subspaces, so the adversary can always learn whether a bit was disturbed ($Z_1 = X \oplus Y$ in the notation of Ref. [21], p. 2) without affecting anything else. After a b-bit block is accepted and both parity strings are announced, she knows whether the block was accepted with all bits correct (probability $p_u = (1-D)^b/p_a$, $p_a = (1-D)^b + D^b$) or all bits wrong, and in the first case her b ancillas are in the product state $\psi_{a_1} \otimes \cdots \otimes \psi_{a_b}$, where the parities fix all a_i relative to the distilled bit a_1.

Per-probe measurement (class H). This is the per-probe adversary of Ref. [7] and the model of Ref. [21]: the Helstrom measurement of each ancilla followed by optimal decoding. The Helstrom measurement distinguishes two equiprobable pure states of overlap c with success probability $\frac{1}{2} + \frac{1}{2}\sqrt{1-c^2}$ [28] (Eq. (3.12) of Ref. [33]), so each ancilla yields a guess of its bit that is correct with probability $\delta_0 = \frac{1}{2} + \frac{1}{2}\sqrt{1-c_\psi^2}$ on undisturbed bits and $\delta_1 = \frac{1}{2} + \frac{1}{2}\sqrt{1-c_\varphi^2}$ on disturbed ones; with the parities the guesses become b votes on a_1. The tally k of votes for 0 is a sufficient statistic, the majority vote is its maximum-likelihood decision with error $\gamma_b(\delta)$, and the adversary's information on the distilled bit is

$$I_E^{\mathrm{ind}}(D,b) = \max_{c_\varphi} \Big[p_u\, J_b(\delta_0) + (1-p_u)\, J_b(\delta_1) \Big], \tag{3}$$

$$J_b(\delta) = 1 - \sum_{k=0}^{b} \Pr(k)\, h\big(\Pr(a_1 = 0 \,|\, k)\big),$$

with $\Pr(a_1 = 0|k) = \delta^k(1-\delta)^{b-k}/[\delta^k(1-\delta)^{b-k} + (1-\delta)^k\delta^{b-k}]$. We stress that $J_b(\delta) \ge 1 - h(\gamma_b(\delta))$ (Fano's inequality; Appendix C.2): the expression built from the majority-vote error, which is what the asymptotic analysis of Ref. [21] needs, is only a lower bound on the adversary's information and must not be used in a key-length formula. At the error rates relevant for key generation the maximum is at $c_\varphi = 1$, where Eq. (3) of Ref. [21] gives $c_\psi = (1-3D)/(1-D)$ (Appendix C.1); close to $D = 1/4$ the maximiser moves below $c_\varphi = 1$ (for $b = 2$ at $D \ge 0.206$, $b = 3$ at 0.222, $b = 5$ at 0.232, $b = 9$ at 0.239 and $b = 15$ at 0.243), and the numerical results below maximise over c_φ throughout. This is the incoherent-attack model of Ref. [21]; its Theorem 1

(Eq. (11) there) states that key agreement by AD is possible iff $D/(1-D) < 2\sqrt{\delta_0(1-\delta_0)} = |c_\psi|$, which for the 4-state protocol is $D < 1/4$ (Eq. (8) there), and it is the individual-attack security condition for the parity-announcing protocol in Ref. [7].

Joint measurement (class J). If the adversary may measure the b ancillas together, but still before the parities are announced, her task is to discriminate the 2^b product states $\psi_{a_1}\otimes\cdots\otimes\psi_{a_b}$ with the post-measurement information c, which partitions them into 2^{b-1} pairs $\{\psi_a,\psi_{\bar a}\}$ [25, 27]. Her success probability is the value of the semidefinite programme

$$P_b^{\mathrm{joint}} = \max_{\{M_f\}}\ 2^{-b}\sum_a \mathrm{Tr}\big[\rho_a E_a\big], \qquad (4)$$
$$E_a = \sum_{f:\, f(c(a))=a_1} M_f,\quad M_f\geq 0,\quad \textstyle\sum_f M_f = \mathbb{1},$$

over POVMs indexed by the functions $f:\{0,1\}^{b-1}\to\{0,1\}$ that assign a guess to every possible parity string; the programme is the standard semidefinite formulation of minimum-error discrimination [29] with the outcome alphabet enlarged to the functions f [25]. It is upper-bounded by the adaptive value $P_b^{\mathrm{ad}} = \frac{1}{2}+\frac{1}{2}\sqrt{1-c_\psi^{2b}}$, the Helstrom probability for the two product states of a known pair, since the overlap of two product states is the product of the single-copy overlaps, c_ψ^b (this is Eq. (17) of Ref. [6] and Eq. (30) of Ref. [7]); this is the attack of Refs. [6, 7], which an adversary with memory implements after the announcement by an adaptive sequence of single-copy measurements [8]. The bound is attained for every b:

Theorem 1. Let ψ_0,ψ_1 be pure states with real overlap c, $|c|<1$, and write $\Psi_a=\psi_{a_1}\otimes\cdots\otimes\psi_{a_b}$. In a basis with $\psi_1 = Z\psi_0$ let $\Psi=\psi_0^{\otimes b}$, $G=Z^{\otimes b}$, $u\propto(1+G)\Psi$, $v\propto(1-G)\Psi$ normalised, and $g_\pm=(u\pm v)/\sqrt{2}$. For $y'\in\{0,1\}^{b-1}$ put $\tilde y=(1\oplus|y'|,\,y')$, let L be the prefix-parity map $(Ly')_j=\bigoplus_{i\leq j}y'_i$, and $P_{y'}=X^{\tilde y}Z^{(0,Ly')}$. Then $\{P_{y'}g_+,P_{y'}g_-\}_{y'}$ is an orthonormal basis, and the projective measurement in it, read after the parities with the rule "$a_1=0$ iff the outcome sign is $(-1)^{y'\cdot(Ly'\oplus c)}$", guesses a_1 with probability $P_b^{\mathrm{ad}}=\frac{1}{2}+\frac{1}{2}\sqrt{1-c^{2b}}$ for every parity class c.

The proof (Appendix A) is short: the basis is a Pauli orbit of the class-0 Helstrom pair, orthonormality follows from $\langle\psi_0|XZ|\psi_0\rangle=0$ for real states together with a property of the prefix-parity map, and the Helstrom projector of every class is a sum of basis projectors because $X^{\tilde y}$ anticommutes with G and commutes with the class translations $Z^{(0,c)}$ up to a sign. In the picture in which the probes carry bit flips, the outcome $(\tilde y,s)$ is a guess of the odd-weight parity $\tilde y\cdot a$, which distinguishes a from $\bar a$ for every class; the post-measurement information only tells the adversary how to translate that parity into a_1. Independently, solving the semidefinite programme for $b=2,3,4$ and $D\in[0.05,0.24]$ gives $P_b^{\mathrm{joint}}=P_b^{\mathrm{ad}}$ to within 10^{-7} (Table II), and the construction of Theorem 1 reproduces P_b^{ad} to 10^{-14} for $b\leq 9$ in our implementation. Hiding the parities therefore costs the adversary nothing in guessing probability; the compatibility of the class-wise optimal measurements that the theorem establishes is the property studied in general by Carmeli, Heinosaari and Toigo [26]. For the adversary's Shannon information we use the fact that for two equiprobable pure states the measurement that maximises the mutual information coincides with the Helstrom measurement, a result of Levitin [32] reproduced in Chapter 3 of Ref. [33] (Eq. (3.263) there gives the accessible information as a function of the overlap), so that $I_E^{\mathrm{joint}}(D,b) = \max_{c_\varphi}[p_u(1-h(1-P_b^{\mathrm{ad}}(c_\psi)))+(1-p_u)(1-h(1-P_b^{\mathrm{ad}}(c_\varphi)))]$ bounds the accessible information of any within-block strategy, adaptive or not; the guessing probability P_b^{ad} bounds the min-entropy version without further assumptions.

Is the Helstrom benchmark the best per-probe measurement? Class W allows any identical single-probe POVM. The probe states span a qubit, so an extremal POVM has at most four rank-one outcomes; we restrict the search to such POVMs (a restriction, since the block information is not convex in the single-probe POVM), and the block information for b identical measurements is then a function of the single-probe outcome statistics (Appendix C.2). Maximising it numerically over real four-outcome POVMs (`perprobe_opt.py`) reproduces the Helstrom-plus-tally values to 10^{-6} in both J_b and the guessing probability for $b=3,5,7,9$ at four values of D in $[0.10,0.23]$ (a local search with twelve restarts, so without a guarantee of global optimality), whereas for even b a tilted measurement that avoids the tie of the majority vote does better in guessing probability (0.8454 against 0.8143 at $b=2$, $D=0.10$; 0.9246 against 0.9093 at $b=4$) and marginally in information. Role-dependent measurements, which do better still at even b, are not available in class W because the roles are fixed by the partition. We therefore use the Helstrom benchmark for odd b only.

What hiding the partition is worth to an adversary with memory. If the adversary can hold n probes but does not know which of them will share a block, the partition joins the parities as post-measurement information. For $b=2$ we solved the corresponding semidefinite programme, with the block containing probe 1 chosen uniformly among the $n-1$ candidates (`hidden_partition_sdp.py`). The optimal success probability is 0.8938 at $n=3$, 0.8913 at $n=4$, 0.8897 at $n=5$, 0.8884 at $n=6$ at $D=0.10$, against the adaptive value 0.8981 and the per-probe benchmark 0.8143 (Table IV). An adversary holding six probes without the partition thus loses less than one percentage point of guessing probability against the adaptive value, recovering more than four fifths of the gap between the best identical per-probe measurement (0.8454) and the adaptive value, and the loss grows slowly with n. Hiding the partition therefore protects only through locality: the

probes of a block must lie in different memory windows, which is what the random partition of Ref. [1] combined with $t_{\mathrm{coh}} \ll T_{\mathrm{run}}$ provides.

The consequence for the asymptotic threshold is immediate (Appendix C.3). Bob's error on an accepted block is $D^b/p_a \approx (D/(1-D))^b$. The adversary's error is $\approx |c_\psi|^b$ for individual measurement (the Chernoff exponent of the majority vote, $2\sqrt{\delta_0(1-\delta_0)} = |c_\psi|$) and $\approx c_\psi^{2b}/4$ for joint measurement (expanding P_b^{ad}). Key agreement by the repetition code requires the adversary's error to decay more slowly than Bob's, the argument of the proof of Theorem 1 of Ref. [21], i.e. $|c_\psi| > D/(1-D)$ in the first case and $c_\psi^2 > D/(1-D)$ in the second. With $c_\psi = (1-3D)/(1-D)$ at $c_\varphi = 1$ these conditions read $D < 1/4$ and $10D^2 - 7D + 1 > 0$, i.e. $D < 1/5$. The second condition is the security condition of Ref. [7] (their Eq. (68) with $d = 2$), which that work derived for an adversary who measures after the announcement; Theorem 1 shows that the same condition governs an adversary who measures before it. The same algebra for the six-state protocol, whose symmetric attack has $c_\varphi = 0$ (Ref. [21], p. 2) and hence $c_\psi = (1-2D)/(1-D)$ by Eq. (3) of that work, gives $5D^2 - 5D + 1 > 0$, i.e. $D < (5-\sqrt{5})/10 = 27.6\%$ [5, 6]. The 25% limit of Ref. [21] is thus the limit for an adversary whose memory is shorter than the spacing between the pulses of a block, not for one whose memory is shorter than the announcement delay.

General Bell-diagonal attacks. The symmetric attack has equal error rates in the two bases. After a random Pauli twirl applied and announced by Alice and Bob, any collective attack on the single-photon rounds is Bell-diagonal [35]: Alice, Bob and the adversary share $\sum_i \sqrt{\lambda_i}|\beta_i\rangle|e_i\rangle$ with the four Bell states β_i, orthonormal e_i, and weights with $\phi_Z = \lambda_3 + \lambda_4$ and $\phi_X = \lambda_2 + \lambda_4$. Conditioned on Alice's and Bob's Z outcomes the adversary holds $\psi_a \propto \sqrt{\lambda_1}e_1 + (-1)^a\sqrt{\lambda_2}e_2$ when the bit was undisturbed and $\varphi_a \propto \sqrt{\lambda_3}e_3 + (-1)^a\sqrt{\lambda_4}e_4$ when it was flipped, so the structure of the symmetric attack is preserved with $c_\psi = (\lambda_1 - \lambda_2)/(1-\phi_Z)$ and $c_\varphi = (\lambda_3 - \lambda_4)/\phi_Z$, and Theorem 1 applies. Eliminating $\lambda_4 = \phi_Z(1-c_\varphi)/2$,

$$c_\psi = \frac{1 - 2\phi_X - \phi_Z c_\varphi}{1-\phi_Z}, \qquad (5)$$
$$0 \le \lambda_4 \le \min(\phi_X, \phi_Z), \quad \lambda_1 = 1 - \phi_Z - \phi_X + \lambda_4 \ge 0,$$

which reduces to $c_\psi = (1-2D-Dc_\varphi)/(1-D)$ at $\phi_X = \phi_Z = D$. The derivation is given in Appendix C.4. The adversary's information $I_E(\phi_Z, \phi_X, b)$ is the maximum over c_φ of the expressions above with this c_ψ; it increases with ϕ_X, because a larger phase error lets her choose a smaller $|c_\psi|$, but it is not monotone in ϕ_Z at every ϕ_X, so the key length below takes the supremum over the true Z error below its estimate.

Hiding the basis as well. A protocol that keys from both bases and delays the sifting announcement beyond t_{coh} confronts the adversary with the basis as additional post-measurement information. Without advantage distillation ($b = 1$) this is the memoryless-adversary setting solved by Bocquet, Alléaume and Leverrier [24], who optimised Eve's interaction and POVM and found thresholds of 15.4% for BB84 (against 14.6% with memory), 17.6% for SARG04 and 20.4% for six-state (against 15.6%); in guessing probability the gain is confined to protocols with more than two bases for the probe states considered here (our semidefinite programme below), while in Shannon information Ref. [24] still finds a gain of 0.8 percentage points for BB84. We extended the semidefinite programme to hidden bases and parities, with Eve's probe states conditioned on Bob's outcome and the accepted-block mixture of all-correct and all-flipped blocks. For BB84 the worst-case attack ($c_\varphi = 1$) gives $P^{\mathrm{joint}} = P^{\mathrm{ad}}$ at $b = 1$ and $b = 2$ for every D; the gap that exists at $c_\varphi = 0$ is irrelevant because the adversary chooses c_φ. For six-state at $b = 1$ we reproduce the qualitative gain of Ref. [24] (e.g. $P^{\mathrm{joint}} = 0.751$ against $P^{\mathrm{ad}} = 0.835$ at $D = 0.168$). But as soon as blocks of two bits are formed the gain collapses: for blocks whose two pulses share an unknown basis, $P^{\mathrm{joint}} = P^{\mathrm{ad}}$ to 10^{-6} (0.888644 against 0.888642 at $D = 0.167$, verified with two solvers). For the realistic blocks with independently chosen hidden bases, where the full programme has 2^{18} elements, we bracket the adversary's success: the relaxation in which only the bases are hidden and the parities are known gives $P^{\mathrm{joint}} \le 0.8878$ at $D = 0.167$ and ≤ 0.9556 at $D = 0.259$, and an explicit 16-outcome projective measurement with both hidden, found by first-order optimisation and therefore possibly short of the true optimum, achieves $P \ge 0.8826$ and 0.9395, against $P^{\mathrm{ad}} = 0.8886$ and 0.9583. Hiding both therefore costs the adversary between 0.0008 and 0.006 in guessing probability at $D = 0.167$ and between 0.003 and 0.019 at $D = 0.259$, already at $b = 2$, against 0.084 at $b = 1$ and $D = 0.168$; in terms of her error probability the ratio to the adaptive value falls from 1.5 at $b = 1$ to at most 1.05 at $b = 2$ and at most 1.04 at $b = 3$ at $D = 0.167$; at $D = 0.259$, close to the six-state limit, our achievable bounds are looser and only give 1.46 and 1.43 at $b = 2, 3$, while the relaxation shows that the true ratio at $b = 2$ is at least 1.06. The construction of Theorem 1 does not extend to hidden bases, because the class translations are then the basis-change unitaries on the adversary's probe rather than Pauli operators, which is why a residual gain survives at finite b. At these block sizes the gain from hiding the bases on top of advantage distillation is therefore small but not zero: the one-way memoryless gain of Ref. [24] (20.4% for six-state) is dominated by advantage distillation without any assumption (27.6%), and adding the timelock to advantage distillation is worth between 0.001 and 0.03 bits of min-entropy per distilled bit at $b = 2$ for the two error rates and the symmetric six-state family considered. The threshold is set by the error exponents of Bob and of the adversary; whether the error ratio stays bounded for all b is open, the trend from $b = 1$

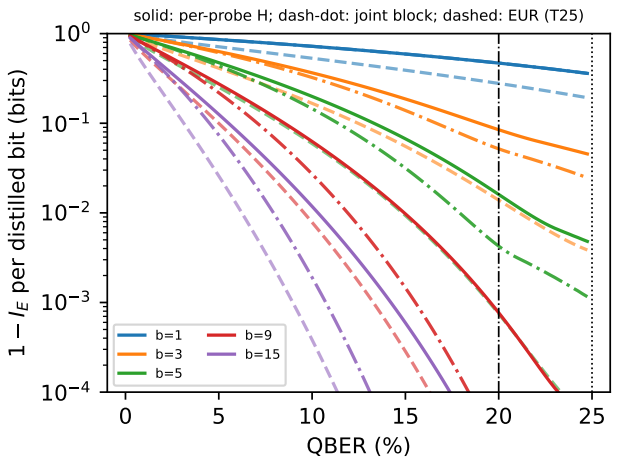


FIG. 1. Adversary's uncertainty per distilled bit, $1 - I_E$, versus QBER for block sizes $b = 1, 3, 5, 9, 15$: Helstrom per-probe benchmark, class H (solid, Eq. (3) with tally decoding), joint block measurement (dash-dotted, adaptive Helstrom value), and the entropic-uncertainty term of Ref. [1] (dashed, same colour). Vertical lines: 20% and 25%.

to $b = 3$ being evidence rather than proof, so we make no claim about the six-state threshold with hidden bases.

Fig. 1 compares $1 - I_E$ in the two models with the entropic-uncertainty term $1 - h(\Phi_{\bar{K}})$ of Ref. [1] at the same D and b.

IV. FINITE-KEY SECURE KEY LENGTH

The protocol is Protocol 1 of Ref. [1] with two modifications: the parity strings c_A, c_B and the reconciliation syndrome are sent under E_k (or after τ_{ann}), and privacy amplification uses the bound below. Let $S^-_{\bar{K}}$ be the lower bound on the number of accepted blocks consisting only of single-photon events (Eq. (18) of Ref. [1]), ϕ^+_Z and ϕ^+_X the upper bounds on the single-photon error rates (Eqs. (17a)–(17b) there), $N_{\bar{K}}$ and $\Phi_{\bar{K}}$ the expected number and error rate of the distilled bits (Eq. (20) there), and $\lambda_{\text{IR}} = fN_{\bar{K}}h(\Phi_{\bar{K}})$ the reconciliation leakage with efficiency f (Sec. 2 of Ref. [1]). Then

$$\begin{aligned}\ell_{\text{QCT}} \le\; & S^-_{\bar{K}}\big[1 - \bar{I}_E\big] \\ & - \Delta_{\text{AEP}}(S^-_{\bar{K}}, \epsilon) - \lambda_{\text{IR}} - \log_2 \frac{2}{\epsilon_{\text{cor}}} \\ & - 4\log_2 \frac{2^{7/4}}{\epsilon_{\text{sec}} - \epsilon^2_{\text{sec}}/8},\end{aligned} \tag{6}$$

with $\bar{I}_E = \sup\{I_E(\phi'_Z, \phi^+_X, b) : 0 < \phi'_Z \le \phi^+_Z\}$ for the individual or the joint model, where $\Delta_{\text{AEP}} = 4\log_2(2+\sqrt{2})\sqrt{S\log_2(2/\epsilon^2)}$ is the finite-size correction of the asymptotic equipartition property (Theorem 9 of Ref. [23], with $\eta \le 2+\sqrt{2}$ for a binary variable and classical side information, valid for $S \ge \frac{8}{5}\log_2(2/\epsilon^2)$), evaluated with smoothing parameter $\epsilon = \epsilon_{\text{sec}}/4$, and the last two terms are the correctness and leftover-hash terms of Eq. (93) of Ref. [1] (Lemma 4 there, with the statistical failure budget $\bar{\epsilon} = \epsilon^2_{\text{sec}}/368$ chosen in Sec. 6 of that work). The only change with respect to Eq. (93) of Ref. [1] is the replacement of the entropic-uncertainty term $S^-_{\bar{K}}[1-h(\Phi^+_{\bar{X}})]$ by $S^-_{\bar{K}}[1-\bar{I}_E]-\Delta_{\text{AEP}}$; Appendix C.5 traces every term to its source. We keep the parameter-estimation and hashing budget of Ref. [1] unchanged and charge the smoothing to an additional $\epsilon_{\text{sec}}/4$, so the total secrecy parameter of the stated numbers is $5\epsilon_{\text{sec}}/4$. A more conservative variant replaces $S^-_{\bar{K}}[1-\bar{I}_E]-\Delta_{\text{AEP}}$ by $-S^-_{\bar{K}}\log_2\bar{p}_{\text{guess}}$ with $\bar{p}_{\text{guess}}$ the same supremum of the adversary's per-block guessing probability ($1-\gamma_b$ or P^{ad}_b, mixed over the two block types), which needs no equipartition step because the min-entropy of a classical bit given classical side information is $-\log_2 p_{\text{guess}}$ (Theorem 1 of Ref. [30]) and is additive over the independent blocks of a collective attack. Both suprema are taken separately over c_φ and over the true Z error, so that each accounting is evaluated at its own worst case; ϕ^+_X can be used directly because both quantities increase with ϕ_X.

Three remarks delimit what this formula is. First, the two accountings play different roles. In the min-entropy accounting with P^{ad}_b the bound is the conditional min-entropy of the distilled bits given the adversary's quantum ancillas and all announcements, because the min-entropy of a product of classical-quantum states is additive (it is a semidefinite programme whose feasible solutions tensorise [30]) and P^{ad}_b is her guessing probability with full quantum memory; this column therefore holds against any measurement, joint across blocks or not, and does not use the timelock at all. It is a collective-attack bound in its own right, and by Theorem 1 the timelock cannot improve it. In the Shannon accounting the term $1 - I^{\text{joint}}_E$ exceeds the Holevo term $1-\chi$ of a quantum-memory adversary, and this excess is the only place where the timelock enters; it is valid for adversaries whose measurement is a product over blocks, since the asymptotic equipartition property is applied to independent classical outcomes, and we do not claim it against joint measurements across blocks, for which accessible information need not be additive. Second, the adversary term is evaluated over all Bell-diagonal collective attacks, which after the twirl of Appendix B are all collective attacks on the untagged rounds; the extension to coherent attacks is discussed in Appendix B and is not included in the numbers, except in Table III, where as a sensitivity analysis we charge the post-selection factor to the smoothing and hashing parameters without completing the reduction. Third, the total secrecy parameter of every number in this paper is $5\epsilon_{\text{sec}}/4$ with $\epsilon_{\text{sec}} = 10^{-9}$, that of Ref. [1] being ϵ_{sec}; rescaling to the same total costs 1.3 bits in the min-entropy accounting and about 10^5 bits in the Shannon accounting through Δ_{AEP}, which moves the thresholds by less than 10^{-6}

percentage points at $N = 10^{14}$ (18.348% at $b = 7$ and 17.659% at $b = 11$ in either normalisation). If A2 fails, the adversary may hold her ancillas until the parities are public, the Shannon replacement is no longer valid, and a timelock-free term must be used: the collective min-entropy term above, or the entropic-uncertainty term of Ref. [1], with which the formula coincides with Eq. (93) of that work because every other term is shared. The key length vanishes for $\phi_Z^+ \geq 1/4$ by construction. Within the model, once the adversary's side information is classical the leftover hash lemma applies in its classical form, the trace distance of the final key from uniform given her classical record is bounded by $5\epsilon_{\rm sec}/4$ conditional on A1 holding during the run, and a trace distance δ from the ideal key bounds the adversary's information on a key of length ℓ by $I(K;E) \leq 2\delta\ell + (1+\delta)h(\delta/(1+\delta))$ through the continuity of conditional entropy (Lemma 2 of Ref. [31]), which is how the accessible-information criterion of Ref. [19] follows from the trace-distance one; the key is everlasting in the sense of Ref. [19] because A1 is needed only then.

V. NUMERICAL RESULTS

We implemented Eqs. (6)–(20), (93) and (94) of Ref. [1] and the formula above in a deterministic Python module, optimising $\{\mu_1, \mu_2, p_{\mu_1}, p_{\mu_2}, p_Z, q_T\}$ by the Nelder–Mead method and locating thresholds by bisection in the misalignment angle for a lossless, noiseless channel ($\eta = 1$, $p_{\rm noise} = 0$, QBER from misalignment only), with $\epsilon_{\rm sec} = 10^{-9}$, $\epsilon_{\rm cor} = 10^{-15}$, $f = 1.2$ and $\mu_3 = 0$, the settings of Ref. [1]. The implementation reproduces the four thresholds reported there: 8.67% and 11.37% at $N = 10^8$ (published 8.7% and 11.3%, $b = 3$) and 9.53% and 17.30% at $N = 10^{14}$ (published 9.5% and 17.3%, $b = 9$).

Table III reports two further sets of numbers. The first charges the post-selection factor of Appendix B, $(N+1)^{15}$, to every term it affects: the leftover-hash term grows by $60\log_2(N+1) \approx 2.8 \times 10^3$ bits, the equipartition term $\Delta_{\rm AEP}$ grows by a factor $\sqrt{1460/64} \approx 4.8$ (about 7×10^7 bits at the optimum), and the parameter-estimation intervals, which we do not shrink, would add a further loss; the min-entropy columns move by less than the resolution of the threshold search (0.05 percentage points), the Shannon columns drop by 0.6–1.5 percentage points. The second replaces $1 - I_E^{\rm joint}$ by the Holevo term $1 - \chi$ of a quantum-memory adversary in the same machinery, which is the timelock-free counterpart of the joint Shannon column.

Asymptotically (single-photon source, $f = 1.2$, evaluated in 60-digit arithmetic because double precision underflows beyond $b \approx 21$), the per-probe threshold rises with b toward 25% (21.9% at $b = 9$, 23.8% at $b = 31$, 24.3% at $b = 61$), whereas the joint Shannon threshold rises to 19.6% at $b = 21$ and 19.9% at $b = 61$, the collective min-entropy threshold to 19.5% at $b = 61$, and the entropic-uncertainty threshold of Ref. [1] to 19.5% at $b = 61$, all three approaching 20% (Fig. 3). At $N = 10^{14}$ the collective min-entropy accounting gives 17.30% at $b = 9$ and 17.66% at $b = 11$, against 17.30% for Ref. [1] at its best $b = 9$; the joint Shannon accounting gives 17.40% at $b = 9$ and 18.35% at $b = 7$, and the Holevo accounting 16.76% at $b = 7$, so the best timelock-dependent certificate exceeds the best timelock-free one by 0.7 percentage points (18.35% at $b = 7$ against 17.66% at $b = 11$, the two in different accountings) and exceeds the Holevo Shannon accounting at the same $b = 7$ by 1.6 percentage points, both for block-local measurements; with the post-selection factor charged the first difference is 0.3 percentage points (17.80% against 17.52% at $b = 11$, Table III). The per-probe benchmark gives 20.66% and 20.81% at $b = 9$ (its best over b is 21.83% at $b = 15$ in the min-entropy accounting), and 20.61% and 19.39% with the post-selection factor charged. The entries of Table I are maximised over all Bell-diagonal attacks with the estimated ϕ_X^+ and ϕ_Z^+; restricting the adversary to the symmetric family $\phi_X = \phi_Z$ raises them by at most 0.16 percentage points at $N = 10^{14}$ and by up to 0.8 percentage points at $N = 10^8$, where the two single-photon error rates differ more through the smaller test sample.

VI. IMPLEMENTATION CONSIDERATIONS

The AD stage adds $b - 1$ exclusive-or operations and one comparison per block, fewer than one operation per raw bit, orders of magnitude below the per-bit cost of iterative LDPC decoding; hardware key-distillation engines [37–39] therefore remain decoder-limited after an AD stage is inserted, whether or not the announcements are delayed. The timelock itself is a buffer of depth $\tau_{\rm ann}$ or an authenticated-encryption wrapper under a post-quantum key-encapsulation mechanism. Given Theorem 1, neither is worth paying for in order to hide the parities from an adversary who knows the partition, since her guessing probability is unchanged and only the Shannon excess of Sec. V is affected; a timelock on the partition is a buffer of the same depth applied to one further message.

VII. DISCUSSION

Three announcements of the protocol can be hidden, and the results above treat them separately. Hiding the parities and the syndrome is settled by Theorem 1: for an adversary who knows which probes form a block, the parity-conditioned discrimination problems are compatible, so the delay does not change her guessing probability, and the collective-attack min-entropy bound, 17.66% at $N = 10^{14}$, is unchanged by the delay. That this bound exceeds the 17.30% of Ref. [1] is a statement about the entropic-uncertainty technique and the attack class (collective here, coherent there), not about the timelock; the only timelock-dependent quantity is the excess of $1 - I_E^{\rm joint}$

TABLE I. Threshold QBER (sifted-key QBER ϕ_K), maximised over Bell-diagonal collective attacks, with the reference values of Ref. [1]. The per-probe columns are the Helstrom benchmark of class H, reported for odd b only (for $b = 2$ it is not the best per-probe measurement, Sec. III); the collective min-entropy column does not depend on the timelock; the joint Shannon column is the timelock-dependent accounting for block-local measurements. Secrecy parameter $5\epsilon_{\rm sec}/4$, no post-selection factor. Replacing the block-count statistic of Ref. [1] by the virtual-i.i.d. Chernoff bound of Ref. [10] changes the $N = 10^{14}$ thresholds by less than 10^{-5}.

N	b	BB84 [T25]	AD [T25]	per-probe H, min-ent	per-probe H, Shannon	collective (joint), min-ent	joint, Shannon
10^8	2	8.67%	10.76%	—	—	9.96%	13.24%
10^8	3		11.37%	13.56%	14.44%	12.01%	13.33%
10^{14}	5	9.53%	16.40%	18.16%	20.18%	15.66%	18.21%
10^{14}	7		17.07%	19.70%	20.91%	16.71%	18.35%
10^{14}	9		17.30%	20.66%	20.81%	17.30%	17.40%
10^{14}	11		16.67%	21.24%	19.20%	17.66%	15.29%
10^{14}	15		13.23%	21.83%	15.45%	17.48%	11.86%

TABLE II. Guessing probability of the distilled bit of an accepted undisturbed block: per-probe Helstrom plus majority vote (for even b ties are broken at random, so $b = 4$ equals $b = 3$; the best identical per-probe measurement does better at even b, Sec. III), the optimal non-adaptive joint measurement (semidefinite programme over $2^{2^{b-1}}$ POVM elements, 256 elements on 16 dimensions at $b = 4$), and the adaptive Helstrom bound. The joint and adaptive columns agree to 10^{-8} for $b = 2, 3$ and to 2×10^{-8} for $b = 4$ at all 21 points computed.

b	D	c_ψ	individual $1-\gamma_b$	joint (SDP)	adaptive $P_b^{\rm ad}$
2	0.10	0.778	0.8143	0.8981	0.8981
2	0.20	0.500	0.9330	0.9841	0.9841
3	0.10	0.778	0.9093	0.9412	0.9412
3	0.20	0.500	0.9871	0.9961	0.9961
3	0.24	0.368	0.9964	0.9994	0.9994
4	0.10	0.778	0.9093	0.9653	0.9653
4	0.20	0.500	0.9871	0.9990	0.9990
4	0.24	0.368	0.9964	0.99992	0.99992

TABLE III. Finite-key thresholds with the post-selection factor charged to the smoothing and hashing parameters (rows marked PS); this is a sensitivity analysis, not a coherent-attack certificate, because the parameter-estimation intervals are not shrunk and the reduction of Appendix B is not completed. The last row is the no-timelock Holevo accounting.

model ($N = 10^{14}$)	$b = 5$	$b = 7$	$b = 9$
joint Shannon (PS)	—	17.80%	15.86%
collective min-ent (PS)	—	16.67%	17.25%
per-probe H Shannon (PS)	—	20.52%	19.39%
per-probe H min-ent (PS)	—	19.70%	20.61%
Holevo Shannon, no timelock	16.40%	16.76%	15.59%

over $1 - \chi$ in the Shannon accounting, 0.7 percentage points between the best values over b for block-local measurements and open against joint measurements across blocks.

Hiding the partition is where the question of Ref. [7] lives. That work distinguishes CAD1, in which Alice announces the positions of b equal bits, from CAD2, the parity-announcing form used here and in Ref. [1]; against an adversary with memory both are secure exactly up to 20%, and against an adversary forced to measure before the announcements Ref. [7] notes that CAD1 is still broken at 20%, because Eve knows in advance that a block carries b copies of one bit [8], whereas CAD2 is held off until the entanglement condition $D < 1/4$. The present work locates the assumption behind the second statement: not that the parities are hidden, which Theorem 1 shows to be irrelevant, but that the adversary cannot bring the probes of a block together before she measures. In Protocol 1 of Ref. [1] the partition is drawn at random after the run, so this holds for any adversary with $t_{\rm coh} \ll T_{\rm run}$ (class W). The per-probe benchmark of Table I is what she concedes under the assumptions of A4 (feed-forward on unrelated probes and correlations across blocks neglected), which we have not derived from the window constraint, and with the proviso that the search over identical per-probe measurements found no improvement only for odd $b \leq 9$. The semidefinite programme of Table IV shows, for up to six held probes, that hiding alone is not enough: an adversary who holds six probes without the partition still reaches within one percentage point of the adaptive value, so at these sizes the protection comes from locality; its large-n behaviour is open. *Open question.* For a window adversary who holds the $w = f_s t_{\rm coh}$ sifted single-photon probes of each memory window, f_s being their rate, and learns the partition, the bases and the parities afterwards, what are the op-

TABLE IV. Hidden block partition as post-measurement information ($b = 2$): optimal success probability of an adversary who holds n probes and learns afterwards which of them is the block mate of probe 1, from the semidefinite programme over $2^{2(n-1)}$ POVM elements, compared with the per-probe Helstrom benchmark, the best identical per-probe measurement of Sec. III, and the adaptive value.

n	D	per-probe H	best identical per-probe	hidden partition, n probes held	adaptive P_2^{ad}
3	0.10	0.8143	0.8454	0.8938	0.8981
3	0.20	0.9330	0.9542	0.9825	0.9841
4	0.10	0.8143	0.8454	0.8913	0.8981
4	0.20	0.9330	0.9542	0.9813	0.9841
5	0.10	0.8143	0.8454	0.8897	0.8981
5	0.20	0.9330	0.9542	0.9804	0.9841
6	0.10	0.8143	0.8454	0.8884	0.8981
6	0.20	0.9330	0.9542	0.9797	0.9841

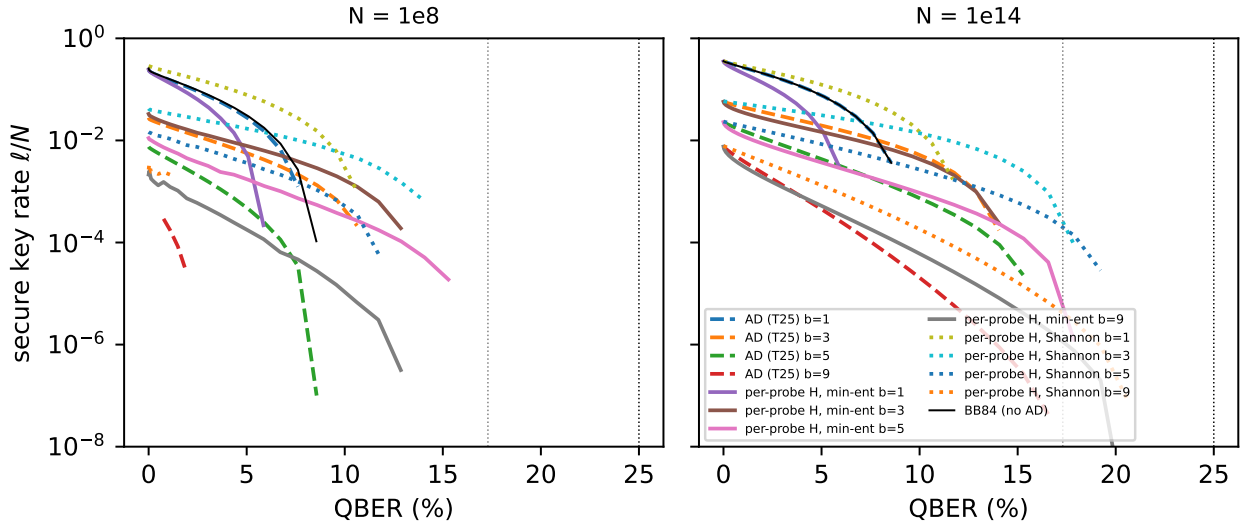


FIG. 2. Secure key rate ℓ/N versus QBER for $N = 10^8$ (left) and $N = 10^{14}$ (right): BB84 one-way (Ref. [1], Eq. (94)), AD (Ref. [1], Eq. (93)) and the Helstrom per-probe benchmark (class H, under the assumptions of A4; not a guarantee against class W) in both entropy accountings, for $b = 1, 3, 5, 9$.

timal guessing probability and accessible information of a distilled bit, as functions of w, b and the Bell-diagonal parameters? A complete statement fixes the number n_K of sifted probes (not the N pulses), their arrival times, whether windows are fixed or each probe has its own lifetime, the order in which bases, partition, acceptance and parities are revealed, the admissible instruments including classical feed-forward, and the attack parameters over which the adversary optimises; the two quantities are separate suprema, and the security of the whole key needs a smooth min-entropy bound rather than a single-bit value. The known points are these. At $w = 1$ the Helstrom benchmark is achievable and, within the identical four-outcome measurements searched in Sec. III, not improved at odd $b \leq 9$. If the adversary holds all n_K probes and learns the partition before she measures, the answer is the adaptive value by Theorem 1; if she learns it only afterwards, the adaptive value is an upper bound that Table IV shows is not attained at $b = 2$ for $n \leq 6$ held probes (0.8884 against 0.8981 at $n = 6$, $D = 0.10$), the large-n limit being open. Table IV is the problem in which the mate of one probe is among n held probes, not the window problem at $w = n$. If the per-probe value persists for $w \ll n_K/b$, a timelock on the partition raises the finite-key threshold of Ref. [1] from 17.3% to about 20.6%; if a window adversary can approach the adaptive value with w far below the block spacing, it raises nothing. Two further steps would be needed for a proof even in the first case: a bound on the fraction of blocks with two probes in one window, which the random partition makes hypergeometric, and the reduction to coherent attacks, which conflicts with block locality because permutation invariance now holds only at the level of blocks (Appendix B).

Hiding the basis is the third lever. For a Z-keyed protocol it changes nothing (Section 2). For protocols that key from both bases the memoryless-adversary thresholds of Ref. [24] apply without AD; our semidefinite programmes

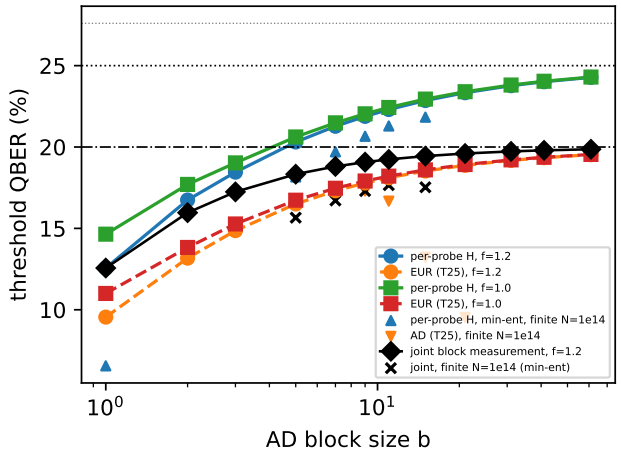


FIG. 3. Threshold QBER versus block size b: finite-key values at $N = 10^{14}$ for Ref. [1], the Helstrom per-probe benchmark (class H, all b) and the joint model, and the asymptotic single-photon thresholds (60-digit arithmetic); horizontal lines at 20% and 25%.

show that for BB84 the worst-case attack is unaffected in guessing probability, that Ref. [24] nevertheless finds 0.8 percentage points in Shannon information, and that for six-state the gain of one-bit blocks shrinks by an order of magnitude at two- and three-bit blocks. Whether it vanishes for large b is open. MUB-QCT [19] is not covered by any of this, since it uses $d + 1$ bases in dimension d and no parity-announcing step; we make no claim about it beyond noting that for qubit protocols the basis and the partition are the only announcements whose hiding is not settled by Theorem 1.

Two observations about the reference formula arose in the course of the work. The bound on the logical phase-error rate of the distilled bits, Eq. (19) of Ref. [1], is proven monotone in the single-photon error rates only on the domain $\phi_X + \phi_Z < 1/2$ (Appendix B of that work), because blocks accepted with all b bits flipped carry no phase error: at $\phi_X = 0.3$ and $b = 7$ the expression falls from 0.499 to 0.25 as ϕ_Z rises from 0.3 to 0.5. An unconstrained optimisation of the protocol parameters leaves the domain by driving the estimated single-photon Z error toward one half through a vanishing test set, and then certifies key at QBERs above the asymptotic limit. Any implementation must enforce $\phi_X^+ + \phi_Z^+ < 1/2$ as an acceptance condition; we do so, and all four thresholds of Ref. [1] lie inside the domain. Separately, the Fano expression $1 - h(\gamma_b)$ built from the majority-vote error, which suffices for the asymptotic threshold of Ref. [21], underestimates a per-probe adversary's Shannon information and must be replaced by the tally information J_b of Sec. III in any finite-key use; the difference is 1.0 percentage points at $N = 10^{14}$, $b = 9$ (21.89% against 20.86% in the symmetric family).

VIII. CONCLUSION

A quantum computational timelock on the post-processing of decoy-state BB84 with advantage distillation can hide the parities, the block partition and the bases. Hiding the parities buys nothing: Theorem 1 gives, for every block size, an explicit measurement made before the parities are known that attains the adaptive Helstrom value, so the adversary's guessing probability is unchanged and the key length in the min-entropy accounting is the collective-attack value, 17.66% at 10^{14} pulses against the 17.30% certified by the entropic-uncertainty bound of Treplin *et al.*; only the Shannon accounting for block-local measurements retains a timelock-dependent excess, 0.7 percentage points between the best values over b. Hiding the partition is the lever that the remark of Bae and Acín requires: an adversary whose memory is shorter than the time between the probes of a block, which the random partition of Treplin *et al.* makes a sizeable fraction of the run, cannot bring them together; under the stated assumption that feed-forward on unrelated probes and correlations across blocks are neglected, we model her by role-blind per-probe measurements, and the Helstrom benchmark, not improved by a numerical search over identical four-outcome measurements at odd block sizes, gives 20.8% (Shannon) and 20.7% (min-entropy) at 10^{14} pulses and $b = 9$, 19.4% and 20.6% with the post-selection factor charged as a sensitivity check, with the 25% limit of Gisin and Wolf as asymptote. Whether an adversary who holds many probes without the partition can do better is open; at six probes she recovers most of the adaptive value, so at these sizes the gain rests on locality. Hiding the basis leaves the BB84 guessing probability unchanged at one- and two-bit blocks and, for six-state, reduces the gain of the memoryless setting by an order of magnitude once two- or three-bit blocks are formed. The key-length formula of Treplin *et al.* must be used within its stated domain, and the per-probe adversary's information must be computed from the vote tally rather than from the majority-vote error.

Data availability. The Python implementation, every result file behind the tables and figures, and the Markdown source of this paper are available at `https://github.com/deeptell-inc/timelocked-ad` and are also provided as Supplemental Material (ancillary files of the arXiv version).

Appendix A: Proof of Theorem 1

Choose the qubit basis $|0\rangle \propto \psi_0 + \psi_1$, $|1\rangle \propto \psi_0 - \psi_1$, so that $\psi_0 = \cos\theta|0\rangle + \sin\theta|1\rangle$, $\psi_1 = Z\psi_0$ and $c = \cos 2\theta$; we assume $|c| < 1$. Then $\Psi_a = Z^a\Psi$ with $\Psi = \psi_0^{\otimes b}$, and $\Psi_\pm = \frac{1}{2}(1 \pm G)\Psi$ have norms $\|\Psi_\pm\|^2 = \frac{1}{2}(1 \pm c^b)$. The class c consists of $Q_c\Psi$ and $GQ_c\Psi$ with $Q_c = Z^{(0,c)}$, and since Q_c commutes with G its Helstrom vectors are

$g_{\pm}^{(c)} = Q_c g_{\pm}$. For any projector E with $Eg_{+}^{(c)} = g_{+}^{(c)}$ and $Eg_{-}^{(c)} = 0$ the success probability on class c is $\frac{1}{2}[\langle Q_c\Psi|E|Q_c\Psi\rangle + \langle GQ_c\Psi|1-E|GQ_c\Psi\rangle] = \frac{1}{2}(\|\Psi_+\| + \|\Psi_-\|)^2 = P_b^{\mathrm{ad}}$, because $Q_c\Psi = \alpha g_{+}^{(c)} + \beta g_{-}^{(c)}$ and $GQ_c\Psi = \alpha g_{-}^{(c)} + \beta g_{+}^{(c)}$ with $\alpha = (\|\Psi_+\| + \|\Psi_-\|)/\sqrt{2}$. It therefore suffices to show that the basis of Theorem 1 is orthonormal (Lemma 1) and that for each class the vectors selected by the rule span $g_{+}^{(c)}$ and are orthogonal to $g_{-}^{(c)}$ (Lemma 2).

For a product state and real ψ_0 one has $\langle\Psi|X^wZ^s|\Psi\rangle = \prod_i \langle\psi_0|X^{w_i}Z^{s_i}|\psi_0\rangle$ with $\langle\psi_0|Z|\psi_0\rangle = c$, $\langle\psi_0|X|\psi_0\rangle = \sin 2\theta \neq 0$ and $\langle\psi_0|XZ|\psi_0\rangle = 0$ (XZ is antisymmetric); hence $\langle\Psi|X^wZ^s|\Psi\rangle = 0$ whenever w and s share a position (and, for $c \neq 0$, only then); Lemma 1 uses only the first direction.

Lemma 1. For $y' \neq z'$ and all signs, $\langle P_{y'}g_s|P_{z'}g_{s'}\rangle = 0$. Indeed $P_{y'}^{\dagger}P_{z'} = \pm X^wZ^{(0,r)}$ with $w = \tilde{y} \oplus \tilde{z} = (|d| \bmod 2, d)$, $d = y' \oplus z' \neq 0$ and $r = Ld$. As $|w|$ is even, $X^wZ^{(0,r)}$ preserves the eigenspaces of G, so the overlap equals $\frac{1}{2}[\langle u|X^wZ^{(0,r)}|u\rangle + ss'\langle v|X^wZ^{(0,r)}|v\rangle]$, and $\langle\Psi_\pm|X^wZ^{(0,r)}|\Psi_\pm\rangle = \frac{1}{2}[\langle\Psi|X^wZ^{(0,r)}|\Psi\rangle \pm \langle\Psi|X^wZ^{(1,\bar{r})}|\Psi\rangle]$. Both expectation values vanish, for every s, s', provided (i) d and Ld share a position and (ii) either $|d|$ is odd or $d \not\subseteq Ld$. For the prefix-parity map, $(Ld)_{i_0} = 1$ at the first position i_0 of d, which gives (i); and if $|d|$ is even, $(Ld)_{i_1} = |d| \bmod 2 = 0$ at the last position i_1 of d, which gives (ii). Within one y', $\langle g_+|g_-\rangle = 0$. The 2^b vectors are therefore orthonormal.

Lemma 2. $\langle P_{y'}g_s|g_{+}^{(c)}\rangle = 0$ unless $s = \epsilon(y', c) := (-1)^{y'\cdot(Ly'\oplus c)}$, and $\langle P_{y'}g_s|g_{-}^{(c)}\rangle = 0$ unless $s = -\epsilon(y', c)$. Using $X^{\tilde{y}}Z^{(0,c)} = (-1)^{y'\cdot c}Z^{(0,c)}X^{\tilde{y}}$, $\langle P_{y'}g_s|Q_cg_+\rangle = (-1)^{y'\cdot c}\langle g_s|Z^{(0,k)}X^{\tilde{y}}|g_+\rangle$ with $k = Ly' \oplus c$. Since $|\tilde{y}|$ is odd, $Z^{(0,k)}X^{\tilde{y}}$ exchanges the eigenspaces of G and only the cross terms survive: the overlap is $(-1)^{y'\cdot c}\frac{1}{2}A[1 + s(-1)^{y'\cdot k}]$ with $A = \langle u|Z^{(0,k)}X^{\tilde{y}}|v\rangle$, where we used $\langle v|Z^{(0,k)}X^{\tilde{y}}|u\rangle = \langle u|X^{\tilde{y}}Z^{(0,k)}|v\rangle = (-1)^{y'\cdot k}A$ for real vectors. This vanishes unless $s = (-1)^{y'\cdot k}$. For $g_{-}^{(c)} = Gg_{+}^{(c)}$, the relations $GX^{\tilde{y}} = -X^{\tilde{y}}G$, $[G, Z^{(0,t)}] = 0$ and $Gg_s = g_{-s}$ give $\langle P_{y'}g_s|g_{-}^{(c)}\rangle = -\langle P_{y'}g_{-s}|g_{+}^{(c)}\rangle$, which vanishes unless $s = -\epsilon(y', c)$.

With $E_c = \sum_{y'} |P_{y'}g_{\epsilon(y',c)}\rangle\langle P_{y'}g_{\epsilon(y',c)}|$, Lemma 2 and the completeness of the basis give $E_cg_{+}^{(c)} = g_{+}^{(c)}$ and $E_cg_{-}^{(c)} = 0$, which proves the theorem. The same construction applies to the all-flipped blocks with φ in place of ψ, and to the six-state probes once the bases are announced.

Appendix B: Twirling, post-selection and block locality

The key length of Sec. IV is evaluated over Bell-diagonal collective attacks on the untagged rounds. (1) Tagging. As in Ref. [1], rounds in which Alice emits more than one photon or Bob's detector does not fire are tagged and the adversary is credited with full knowledge of them; the decoy-state bounds $S_{\bar{K}}^{-}$ and ϕ_Z^{+} are statements about Bob's click record and hold against any attack. (2) Symmetrisation. Alice and Bob may apply a random Pauli twirl to the untagged rounds and announce it at once; it commutes with their measurements, costs no key and reduces any collective attack to a Bell-diagonal one with the observed error rates ϕ_X, ϕ_Z [35, 36], whose purification has the structure of Sec. III with $c_\psi = (1 - 2\phi_X - \phi_Z c_\varphi)/(1 - \phi_Z)$ (Appendix C.4); an adversary holding the purification is at least as strong as one who must measure early, so the reduction is valid in the security direction. (3) Coherent attacks. For a permutation-invariant protocol the post-selection technique [34] lifts security against every i.i.d. attack with parameter ϵ to security against arbitrary attacks with parameter $\epsilon(n+1)^{d^2-1}$, $d = 4$ for a pair of qubits. Two points prevent us from applying it here as more than a comment. First, the factor must be charged to every parameter of the i.i.d. analysis that it affects: the leftover-hash term grows by $4\log_2(n+1)$[15] bits and, in the Shannon accounting, the smoothing parameter of the equipartition term shrinks by the same factor, which is why Table III reports losses of up to 1.5 percentage points rather than the few hundred bits that the hashing term alone would suggest. Second, the block structure breaks round-wise permutation invariance for the classes J and W, since which rounds share a block matters; the technique then applies at the level of blocks, with local dimension 4^b and exponent $16^b - 1$, which is useless for $b \geq 2$. For the collective min-entropy column, which does not depend on the timelock, the entropic-uncertainty proof of Ref. [1] already covers coherent attacks and the question is only whether the sharper collective bound extends; for the timelock-dependent columns, a proof against coherent attacks would need a tool compatible with block locality, such as entropy accumulation, and we leave it open.

Appendix C: Derivations of the expressions in the main text

This appendix collects the algebra behind the expressions of the main text and attributes every step that is not ours to its source, so that the reader can follow the chain from the cited equations to the numbers in Table I.

C.1. Per-ancilla quantities. Eqs. (1)–(2) of Ref. [21] define the symmetric attack, and Eq. (3) there expresses the fidelity through the overlaps, $F = (1 + c_\varphi)/(2 - c_\psi + c_\varphi)$. Since $F = 1 - D$ by definition of the disturbance, solving for c_ψ gives

$$c_\psi = 2 + c_\varphi - \frac{1 + c_\varphi}{1 - D},$$

which is $c_\psi = (1-3D)/(1-D)$ at $c_\varphi = 1$ (the BB84 worst case at the error rates of interest, Sec. III) and $c_\psi = (1-2D)/(1-D)$ at $c_\varphi = 0$, the value that Ref. [21] (p. 2) attributes to the six-state attack. The Helstrom measurement for two equiprobable pure states with overlap c succeeds with probability $\frac{1}{2}+\frac{1}{2}\sqrt{1-c^2}$ [28], rederived as Eq. (3.12) of Ref. [33]; applied to $\{\psi_\uparrow,\psi_\downarrow\}$ and $\{\varphi_\uparrow,\varphi_\downarrow\}$ it gives δ_0 and δ_1. The identity $4\delta_0(1-\delta_0) = 1-(1-c_\psi^2) = c_\psi^2$ turns the condition $D/(1-D) < 2\sqrt{\delta_0(1-\delta_0)}$ of Theorem 1 (Eq. (11)) of Ref. [21] into $|c_\psi| > D/(1-D)$; with $c_\psi = (1-3D)/(1-D)$ this is $1-3D > D$, i.e. $D < 1/4$, which is Eq. (8) of Ref. [21]. That the adversary learns whether a bit was disturbed is the statement $Z_1 = X \oplus Y$ of Ref. [21] (p. 2), a consequence of the orthogonality of the ψ and φ subspaces.

C.2. Majority vote, tally, and general per-probe measurements. After the parities c are announced, the adversary's b per-ancilla guesses $\hat{a}_i$ become b guesses $\hat{a}_i \oplus c_{i-1}$ of the distilled bit a_1, each correct with probability δ (δ_0 or δ_1 according to the block type, which she knows from Z_1) and independent given a_1 because the ancillas are in a product state. Let k be the number of guesses equal to 0. Under $a_1 = 0$ the likelihood of k is $\binom{b}{k}\delta^k(1-\delta)^{b-k}$ and under $a_1 = 1$ it is $\binom{b}{k}(1-\delta)^k\delta^{b-k}$; their ratio, $(\delta/(1-\delta))^{2k-b}$, depends on the guesses only through k, so the tally is a sufficient statistic and the maximum-likelihood decision is the majority vote, with error

$$\gamma_b(\delta) = \sum_{k>b/2}\binom{b}{k}(1-\delta)^k\delta^{b-k} + \tfrac{1}{2}\binom{b}{b/2}\big[\delta(1-\delta)\big]^{b/2},$$

where the last term is present only for even b and accounts for ties broken at random. The adversary's information is the mutual information between a_1 and the sufficient statistic, $J_b(\delta) = 1 - \sum_k \Pr(k)\, h(\Pr(a_1 = 0|k))$ with the Bayes posterior $\Pr(a_1 = 0|k) = \delta^k(1-\delta)^{b-k}/[\delta^k(1-\delta)^{b-k} + (1-\delta)^k\delta^{b-k}]$, as stated in Sec. III. Fano's inequality for a binary variable, $H(a_1|k) \le h(\Pr[\text{error}])$, gives $J_b \ge 1 - h(\gamma_b)$, which is why the majority-vote expression underestimates the adversary. The Chernoff bound for the binomial tail gives $\gamma_b \le [2\sqrt{\delta(1-\delta)}]^b = |c_\psi|^b$ with the exponent tight, which is the decay rate used in the proof of Theorem 1 of Ref. [21]. For a general identical single-probe POVM $\{E_k\}$ with outcome probabilities $p_k(a) = \langle\psi_a|E_k|\psi_a\rangle$ the same reasoning applies with the pair (k_i, c_i) as the per-probe datum: given a_1, each pair is drawn independently from $\Pr(k, c|a_1) = \frac{1}{2}p_k(a_1 \oplus c)$, so the histogram of the $2K$ categories is a sufficient statistic and both J_b and the guessing probability are explicit multinomial sums; `perprobe_opt.py` maximises them over real rank-one four-outcome POVMs in the Naimark parametrisation $E_k = V^{\mathrm{T}}|m_k\rangle\langle m_k|V$ with $\{m_k\}$ an orthonormal basis of $\mathbb{R}^4$.

C.3. Asymptotic conditions. An accepted block is either all-correct, with probability $(1-D)^b$, or all-flipped, with probability D^b (Eq. (20) of Ref. [1] with $\phi_K = D$), so Bob's error on the distilled bit is $\Phi_{\bar{K}} = D^b/[(1-D)^b + D^b] \approx (D/(1-D))^b$. For the adversary, the individual model gives the majority-vote error $\approx |c_\psi|^b$ of C.2, and the joint model gives $1 - P_b^{\mathrm{ad}} = \frac{1}{2} - \frac{1}{2}\sqrt{1-c_\psi^{2b}} \approx c_\psi^{2b}/4$. The repetition code yields a positive rate iff the adversary's error on the distilled bit decays more slowly than Bob's as $b \to \infty$ (the comparison of exponents in the proof of Theorem 1 of Ref. [21], which relies on the classical results cited there); this gives $|c_\psi| > D/(1-D)$ and $c_\psi^2 > D/(1-D)$ respectively. In the joint case with $c_\psi = (1-3D)/(1-D)$ the condition is $(1-3D)^2 > D(1-D)$, i.e. $10D^2 - 7D + 1 > 0$, whose roots are $D = 1/5$ and $D = 1/2$, so $D < 1/5$. For the six-state protocol, $c_\psi = (1-2D)/(1-D)$ gives $(1-2D)^2 > D(1-D)$, i.e. $5D^2 - 5D + 1 > 0$, with roots $(5 \pm \sqrt{5})/10$, so $D < (5-\sqrt{5})/10$. The condition $c_\psi^2 > D/(1-D)$ is Eq. (68) of Ref. [7] specialised to $d = 2$, where it was derived for an adversary who measures after the announcement; Ref. [6] obtains the six-state value in its Eq. (23).

C.4. Bell-diagonal attacks. After a random Pauli twirl applied by Alice and Bob to their qubits the two-qubit state is Bell-diagonal, $\rho_{AB} = \sum_{i=1}^4 \lambda_i|\beta_i\rangle\langle\beta_i|$ with $\beta_{1,2} = (|00\rangle \pm |11\rangle)/\sqrt{2}$ and $\beta_{3,4} = (|01\rangle \pm |10\rangle)/\sqrt{2}$, and the attack is characterised by the weights [35, 36]; Eq. (7) of Ref. [21] is the special case of the symmetric attack. A Z-basis error occurs for β_3, β_4 and an X-basis error (a phase flip) for β_2, β_4, so $\phi_Z = \lambda_3 + \lambda_4$ and $\phi_X = \lambda_2 + \lambda_4$. The adversary holds the purification $\sum_i \sqrt{\lambda_i}|\beta_i\rangle|e_i\rangle$ with orthonormal e_i. Projecting Alice and Bob on $|a\rangle|y\rangle$ in the Z basis gives, up to the factor $1/\sqrt{2}$, the unnormalised adversary states $\sqrt{\lambda_1}e_1 + (-1)^a\sqrt{\lambda_2}e_2$ for $y = a$ and $\sqrt{\lambda_3}e_3 + (-1)^a\sqrt{\lambda_4}e_4$ for $y \ne a$. Normalising, $\psi_a \propto \sqrt{\lambda_1}e_1 + (-1)^a\sqrt{\lambda_2}e_2$ and $\varphi_a \propto \sqrt{\lambda_3}e_3 + (-1)^a\sqrt{\lambda_4}e_4$, with overlaps $c_\psi = (\lambda_1 - \lambda_2)/(\lambda_1+\lambda_2) = (\lambda_1-\lambda_2)/(1-\phi_Z)$ and $c_\varphi = (\lambda_3-\lambda_4)/\phi_Z$, and $\psi_a \perp \varphi_{a'}$ because the e_i are orthonormal; this is the structure of Sec. III. Writing $\lambda_4 = \phi_Z(1-c_\varphi)/2$, $\lambda_3 = \phi_Z - \lambda_4$, $\lambda_2 = \phi_X - \lambda_4$ and $\lambda_1 = 1 - \phi_Z - \phi_X + \lambda_4$ gives $\lambda_1 - \lambda_2 = 1 - \phi_Z - 2\phi_X + 2\lambda_4 = 1 - 2\phi_X - \phi_Z c_\varphi$, hence the expression for c_ψ quoted in Sec. III; the constraints $\lambda_i \ge 0$ are those stated there. At $\phi_X = \phi_Z = D$ the expression reduces to $c_\psi = (1 - 2D - Dc_\varphi)/(1-D)$, which coincides with C.1.

C.5. The key-length formula. Eq. (93) of Ref. [1] reads $\ell_{\mathrm{AD}} \le S^-_{\bar{K}}[1 - h(\Phi^+_{\bar{X}})] - \lambda_{\mathrm{IR}} - \log_2(2/\epsilon_{\mathrm{cor}}) - 4\log_2\big(2^{7/4}/(\epsilon_{\mathrm{sec}} - \epsilon_{\mathrm{sec}}^2/8)\big)$ when the acceptance thresholds are set to the computed bounds (Eq. (21) there). Its first term is the entropic-uncertainty lower bound on the smooth min-entropy of the distilled bits of the single-photon blocks, $S^-_{\bar{K}}$ being Eq. (18) and $\Phi^+_{\bar{X}}$ Eq. (19) of Ref. [1]; the second is the reconciliation leakage; the third is the error-verification cost (Theorem 2 of the reference cited as [42] in Ref. [1]); the fourth is the leftover-hash-lemma term (Lemma 4 of Ref. [1]) after the statistical failure probabilities are set to $\bar{\epsilon} = \epsilon_{\mathrm{sec}}^2/368$ (Sec. 6 of Ref. [1]). Under the timelock the adversary's side information on each single-photon block is

classical, so we replace the first term by a bound on the smooth min-entropy of $S^{-}_{\bar{K}}$ i.i.d. classical bits given classical side information. The fully quantum asymptotic equipartition property, Theorem 9 (Eq. (30)) of Ref. [23], gives $H^{\epsilon}_{\min}(A^S|E^S) \geq S\,H(A|E) - \sqrt{S}\,\delta(\epsilon,\eta)$ with $\delta(\epsilon,\eta) = 4\log_2\eta\sqrt{\log_2(2/\epsilon^2)}$ for $S \geq \frac{8}{5}\log_2(2/\epsilon^2)$, where η is the convergence parameter of Definition 4 (Eq. (24)) of Ref. [23], bounded there by $\sqrt{2^{-H_{\min}(A|E)}} + \sqrt{2^{H_{\max}(A|E)}} + 1$. For a binary A with classical E one has $H_{\min}(A|E) \geq 0$ and $H_{\max}(A|E) \leq 1$, hence $\eta \leq 2 + \sqrt{2}$, which is the constant in Δ_{AEP}. With a_1 uniform, $H(A|E) = 1 - I_E$, and the worst case over the attack parameters and the true error rates below their estimates gives $\bar{I}_E$. The smoothing parameter $\epsilon = \epsilon_{\mathrm{sec}}/4$ is charged in addition to the budget of Ref. [1], giving the total $5\epsilon_{\mathrm{sec}}/4$ stated in Sec. IV. The min-entropy variant uses instead $H_{\min}(A|E) = -\log_2 p_{\mathrm{guess}}$ for classical side information (Theorem 1 of Ref. [30]), additive over the independent blocks of a collective attack, so that no smoothing step and no Δ_{AEP} term are needed. The leftover hash lemma (Lemma 4 of Ref. [1]), applied to the smoothed bound in the Shannon accounting and to the unsmoothed min-entropy in the min-entropy accounting, then yields a key whose trace distance from uniform, given the adversary's record, is at most $5\epsilon_{\mathrm{sec}}/4$ in the first case and ϵ_{sec} in the second (we quote the common upper bound $5\epsilon_{\mathrm{sec}}/4$ throughout); the conversion of this trace distance into a bound on the adversary's information is the continuity bound quoted in Sec. IV, not an identity between the two criteria.

[1] J. Treplin, P. Kleinpaß, D. Orsucci, Finite size analysis of decoy-state BB84 with advantage distillation, arXiv:2511.21665 (2025).

[2] C. C. W. Lim, M. Curty, N. Walenta, F. Xu, H. Zbinden, Concise security bounds for practical decoy-state quantum key distribution, Phys. Rev. A 89, 022307 (2014).

[3] U. M. Maurer, Secret key agreement by public discussion from common information, IEEE Trans. Inf. Theory 39, 733 (1993).

[4] D. Gottesman, H.-K. Lo, Proof of security of QKD with two-way classical communications, IEEE Trans. Inf. Theory 49, 457 (2003).

[5] H. F. Chau, Practical scheme to share a secret key through a quantum channel with a 27.6% bit error rate, Phys. Rev. A 66, 060302 (2002).

[6] A. Acín, J. Bae, E. Bagan, M. Baig, Ll. Masanes, R. Muñoz-Tapia, Secrecy properties of quantum channels, Phys. Rev. A 73, 012327 (2006); arXiv:quant-ph/0411092.

[7] J. Bae, A. Acín, Key distillation from quantum channels using two-way communication protocols, Phys. Rev. A 75, 012334 (2007); arXiv:quant-ph/0610048.

[8] A. Acín, E. Bagan, M. Baig, Ll. Masanes, R. Muñoz-Tapia, Multiple-copy two-state discrimination with individual measurements, Phys. Rev. A 71, 032338 (2005).

[9] R. Bhavsar, J. Moon, J. Bae, Entropy bounds via hypothesis testing and its applications to two-way key distillation in quantum cryptography, arXiv:2602.05870 (2026).

[10] C. Jiang, X.-L. Hu, Z.-W. Yu, X.-B. Wang, Composable security for practical QKD with two way classical communication, arXiv:2102.00739 (2021).

[11] I. B. Damgård, S. Fehr, L. Salvail, C. Schaffner, Cryptography in the bounded quantum-storage model, Proc. 46th IEEE FOCS, 449 (2005).

[12] S. Wehner, C. Schaffner, B. M. Terhal, Cryptography from noisy storage, Phys. Rev. Lett. 100, 220502 (2008).

[13] R. König, S. Wehner, J. Wullschleger, Unconditional security from noisy quantum storage, IEEE Trans. Inf. Theory 58, 1962 (2012).

[14] S. Watanabe, R. Matsumoto, T. Uyematsu, Security of QKD protocol with two-way classical communication assisted by one-time pad encryption, arXiv:quant-ph/0608030 (2006).

[15] V. A. Pastushenko, D. A. Kronberg, Improving the performance of quantum cryptography by using the encryption of the error correction data, Entropy 25, 956 (2023).

[16] M. Tomamichel, A. Leverrier, A largely self-contained and complete security proof for QKD, Quantum 1, 14 (2017).

[17] R. Novak, Security of QKD with one-time-pad-protected error correction and its performance benefits, Entropy 27, 1032 (2025).

[18] C.-H. F. Fung, X. Ma, H. F. Chau, Q.-Y. Cai, Quantum key distribution with delayed privacy amplification and its application to the security proof of a two-way deterministic protocol, Phys. Rev. A 85, 032308 (2012).

[19] N. Vyas, R. Alléaume, Everlasting secure key agreement with performance beyond QKD in a quantum computational hybrid security model, arXiv:2004.10173 (2020).

[20] F. Mazzoncini, B. Bauer, P. Brown, R. Alléaume, Hybrid quantum cryptography from communication complexity, Quantum 9, 1862 (2025).

[21] N. Gisin, S. Wolf, Quantum cryptography on noisy channels: quantum versus classical key-agreement protocols, Phys. Rev. Lett. 83, 4200 (1999).

[22] C. A. Fuchs, N. Gisin, R. B. Griffiths, C.-S. Niu, A. Peres, Optimal eavesdropping in quantum cryptography. I, Phys. Rev. A 56, 1163 (1997).

[23] M. Tomamichel, R. Colbeck, R. Renner, A fully quantum asymptotic equipartition property, IEEE Trans. Inf. Theory 55, 5840 (2009).

[24] A. Bocquet, R. Alléaume, A. Leverrier, Optimal eavesdropping on QKD without quantum memory, J. Phys. A 45, 025305 (2012); arXiv:1106.0329.

[25] M. A. Ballester, S. Wehner, A. Winter, State discrimination with post-measurement information, IEEE Trans. Inf. Theory 54, 4183 (2008).

[26] C. Carmeli, T. Heinosaari, A. Toigo, State discrimination with postmeasurement information and incompatibility of quantum measurements, Phys. Rev. A 98, 012126 (2018).

[27] D. Gopal, S. Wehner, Using post-measurement information in state discrimination, Phys. Rev. A 82, 022326

(2010).
[28] C. W. Helstrom, Quantum Detection and Estimation Theory (Academic Press, New York, 1976).
[29] Y. C. Eldar, A. Megretski, G. C. Verghese, Designing optimal quantum detectors via semidefinite programming, IEEE Trans. Inf. Theory 49, 1007 (2003).
[30] R. König, R. Renner, C. Schaffner, The operational meaning of min- and max-entropy, IEEE Trans. Inf. Theory 55, 4337 (2009).
[31] A. Winter, Tight uniform continuity bounds for quantum entropies: conditional entropy, relative entropy distance and energy constraints, Commun. Math. Phys. 347, 291 (2016).
[32] L. B. Levitin, Optimal quantum measurements for two pure and mixed states, in Quantum Communications and Measurement, edited by V. P. Belavkin, O. Hirota, R. L. Hudson (Plenum, New York, 1995), p. 439.
[33] C. A. Fuchs, Distinguishability and accessible information in quantum theory, PhD thesis, University of New Mexico (1996); arXiv:quant-ph/9601020.
[34] M. Christandl, R. König, R. Renner, Postselection technique for quantum channels with applications to quantum cryptography, Phys. Rev. Lett. 102, 020504 (2009).
[35] B. Kraus, N. Gisin, R. Renner, Lower and upper bounds on the secret-key rate for quantum key distribution protocols using one-way classical communication, Phys. Rev. Lett. 95, 080501 (2005).
[36] R. Renner, N. Gisin, B. Kraus, An information-theoretic security proof for QKD protocols, Phys. Rev. A 72, 012332 (2005).
[37] N. Walenta et al., A fast and versatile QKD system with hardware key distillation and wavelength multiplexing, New J. Phys. 16, 013047 (2014).
[38] J. Constantin et al., An FPGA-based 4 Mbps secret key distillation engine for QKD systems, J. Signal Process. Syst. 86, 1 (2017).
[39] X. Cheng, H.-K. Mao, H.-W. Xu, Q. Li, Large-scale FPGA-based privacy amplification exceeding 10^8 bits for QKD, arXiv:2503.09331 (2025).